

# SS-ZG548: ADVANCED DATA MINING

## Lecture-12: Clustering on Data Stream, Big Data



**Dr. Kamlesh Tiwari**

Assistant Professor

Department of Computer Science and Information Systems Engineering,  
BITS Pilani, Rajasthan-333031 INDIA

March 10, 2018

(WILP @ BITS-Pilani Jan-Apr 2018)

# Clustering over Evolving Data Stream (DenStream)

## Clustering over Evolving Data Stream

With limited memory and one-pass constraint one want to determine arbitrary number of clusters of arbitrary shape by efficiently handling outliers.

### Use DenStream <sup>1</sup>

- Damped window model: Weights with time  $t$  are  $f(t) = 2^{-\lambda \cdot t}$ . (other models landmark and sliding window)
- core-micro-cluster, potential-micro-cluster and outlier-micro-cluster structures
- Guarantees the precision of the weights of the micro-clusters

---

<sup>1</sup>Cao, Feng and Ester, Martin and Qian, Weining and Zhou, Aoying, “Density-Based Clustering over an Evolving Data Stream with Noise.”, in International Conference on Data Mining, pages 328–339, vol 6, SIAM, 2006

# Clustering over Evolving Data Stream


- Beside limited memory and one-pass constraints, we require
  - ▶ No assumption on the number of clusters,
  - ▶ Discovery of clusters with arbitrary shape and
  - ▶ Ability to handle outliers
- DenStream<sup>2</sup>, is a new approach for discovering clusters in an evolving data stream.
- Uses core-micro-cluster to summarize the clusters
- Along with potential core-micro-cluster and outlier micro-cluster structures
- Designed a pruning strategy that guarantees the precision of the weights of the micro-clusters
- User damped window model. Weights with time  $t$  are  $f(t) = 2^{-\lambda \cdot t}$ . (other models landmark and sliding window)

---

<sup>2</sup>Cao, Feng and Ester, Martin and Qian, Weining and Zhou, Aoying, “Density-Based Clustering over an Evolving Data Stream with Noise.”, in International Conference on Data Mining, pages 328–339, vol 6, SIAM, 2006

# Clustering over Evolving Data Stream

- **Definition (core object)** It is an object in whose  $\epsilon$  neighborhood the overall weight of data points is at least  $\mu$
- **Definition (density-area)** A density area is defined as the union of the  $\epsilon$  neighborhoods of core objects
- **Definition (core-micro-cluster)** At time  $t$  it is defined as  $CMC(w, c, r)$  for a group of close points  $p_{i_1}, p_{i_2}, \dots, p_{i_n}$  with time stamp  $T_{i_1}, T_{i_2}, \dots, T_{i_n}$ .  $w = \sum_{j=1}^n f(t - T_{i_j})$ ,  $w \geq \mu$  is the weight.  
 $c = \frac{\sum_{j=1}^n f(t - T_{i_j}) p_{i_j}}{w}$  is the center,  $r = \frac{\sum_{j=1}^n f(t - T_{i_j}) \text{dist}(p_{i_j}, c)}{w}$ ,  $r \leq \epsilon$  is the radius, where  $\text{dist}(p_{i_j}, c)$  denotes the Euclidean distance between point  $p_{i_j}$  and center  $c$
- When a clustering request arrives, each c-micro-cluster will be labeled to get the final result.



# Clustering over Evolving Data Stream

**Definition (potential-micro-cluster)** A potential-micro-cluster (or p-micro-cluster) at time  $t$  for a group of close points  $p_{i_1}, p_{i_2}, \dots, p_{i_n}$  with time stamp  $T_{i_1}, T_{i_2}, \dots, T_{i_n}$  is defined as  $\{\overline{CF^1}, \overline{CF^2}, w\}$ .  
 $w = \sum_{j=1}^n f(t - T_{i_j})$ ,  $w \geq \beta\mu$  is the weight.  $\beta$ ,  $0 < \beta \leq 1$ , is the parameter to determine the threshold of outlier relative to c-micro-clusters.  
 $\overline{CF^1} = \sum_{j=1}^n f(t - T_{i_j})p_{i_j}$  is the weighted linear sum of the points.  
 $\overline{CF^2} = \sum_{j=1}^n f(t - T_{i_j})p_{i_j}^2$  is the weighted squared linear sum of the points. Centre of p-micro-cluster is  $c = \frac{\overline{CF^1}}{w}$  and the radius of a p-micro-cluster  $r = \sqrt{\frac{\overline{CF^2}}{w} + (\frac{\overline{CF^1}}{w})^2}$ ,  $r \leq \epsilon$

# Clustering over Evolving Data Stream

**Definition (outlier-micro-cluster)** An outlier-micro-cluster (or o-micro-cluster) at time  $t$  for a group of close points  $p_{i_1}, p_{i_2}, \dots, p_{i_n}$  with time stamp  $T_{i_1}, T_{i_2}, \dots, T_{i_n}$  is defined as  $\{\overline{CF^1}, \overline{CF^2}, w, t_0\}$ . The definition of  $w, \overline{CF^1}, \overline{CF^2}$  center and radius are the same as p-micro-cluster.  $t_0 = T_{i_1}$ , denotes the creation time of o-micro-cluster which is used to define the life span of o-micro-cluster. However  $w < \beta\mu$ .

**Note:** p-micro-cluster and o-micro-cluster can be maintained incrementally.

**Clustering Algorithm has two parts:**

- Online part of micro-cluster maintenance
- Offline part of generation of final clusters, on demand of user

# Merging of P

p-micro-clusters and o-micro-clusters are maintained in an online way.

---

**Algorithm 1 Merging ( $p$ )**

---

```
1: Try to merge  $p$  into its nearest p-micro-cluster  $c_p$ ;  
2: if  $r_p$  (the new radius of  $c_p$ )  $\leq \epsilon$  then  
3:   Merge  $p$  into  $c_p$ ;  
4: else  
5:   Try to merge  $p$  into its nearest o-micro-cluster  $c_o$ ;  
6:   if  $r_o$  (the new radius of  $c_o$ )  $\leq \epsilon$  then  
7:     Merge  $p$  into  $c_o$ ;  
8:     if  $w$  (the new weight of  $c_o$ )  $> \beta\mu$  then  
9:       Remove  $c_o$  from outlier-buffer and create a  
       new p-micro-cluster by  $c_o$ ;  
10:    end if  
11:   else  
12:     Create a new o-micro-cluster by  $p$  and insert it  
     into the outlier-buffer;  
13:   end if  
14: end if
```

---

# DenStream Algorithm

---

**Algorithm 2 DenStream** ( $DS, \epsilon, \beta, \mu, \lambda$ )

---

```
1:  $T_p = \lceil \frac{1}{\lambda} \log(\frac{\beta\mu}{\beta\mu-1}) \rceil$ ;
2: Get the next point  $p$  at current time  $t$  from data
   stream  $DS$ ;
3: Merging( $p$ );
4: if ( $t \bmod T_p = 0$ ) then
5:   for each p-micro-cluster  $c_p$  do
6:     if  $w_p$  (the weight of  $c_p$ )  $< \beta\mu$  then
7:       Delete  $c_p$ ;
8:     end if
9:   end for
10:  for each o-micro-cluster  $c_o$  do
11:     $\xi = \frac{2^{-\lambda(t-t_o+T_p)}-1}{2^{-\lambda T_p}-1}$ ;
12:    if  $w_o$  (the weight of  $c_o$ )  $< \xi$  then
13:      Delete  $c_o$ ;
14:    end if
15:  end for
16: end if
17: if a clustering request arrives then
18:   Generating clusters;
19: end if
```

---



# Performance




Figure 4: Synthetic data sets




Figure 5: Clustering on DS1, DS2 and DS3

# Performance




Figure 6: Clustering on the evolving data stream EDS

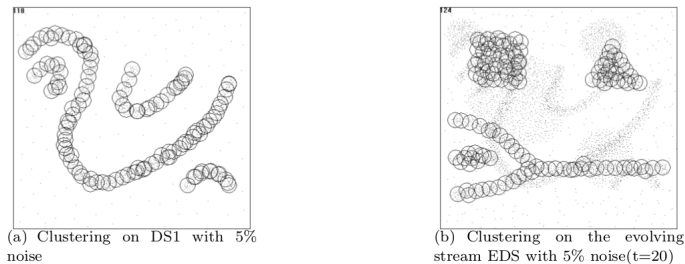


Figure 7: Clustering on data streams with noise

# Performance




Figure 8: Clustering quality(EDS data stream, horizon=2, stream speed=2000)




Figure 9: Clustering quality(EDS data stream, horizon=10, stream speed=1000)




Figure 10: Clustering quality(Network Intrusion dataset, horizon=1, stream speed=1000)




Figure 11: Clustering quality(Network Intrusion dataset, horizon=5, stream speed=1000)

# Performance




Figure 12: Clustering quality(EDS data stream with 1% noise, horizon=2, stream speed=2000)




Figure 13: Clustering quality(EDS data stream with 5% noise, horizon=10, stream speed=1000)




Figure 14: Execution time vs. length of stream(Network Intrusion data set)




Figure 15: Execution time vs. length of stream(Charitable Donation data set)

# Performance




Figure 16: Execution time vs. number of clusters




Figure 18: Memory usage




Figure 17: Execution time vs. dimensionality


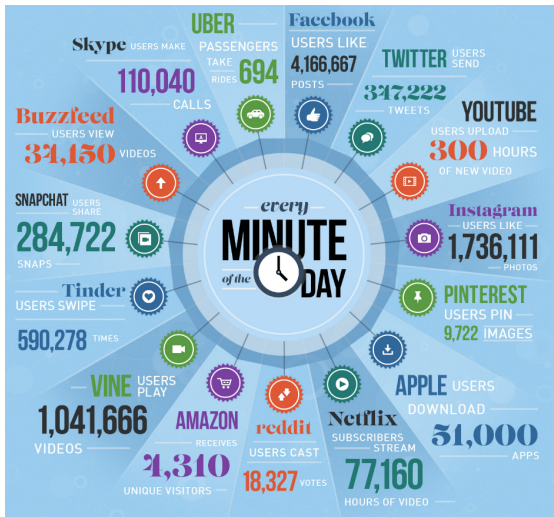


Figure 19: Clustering quality vs. decay factor  $\lambda$

# Big Data

## What is big data?<sup>3</sup>

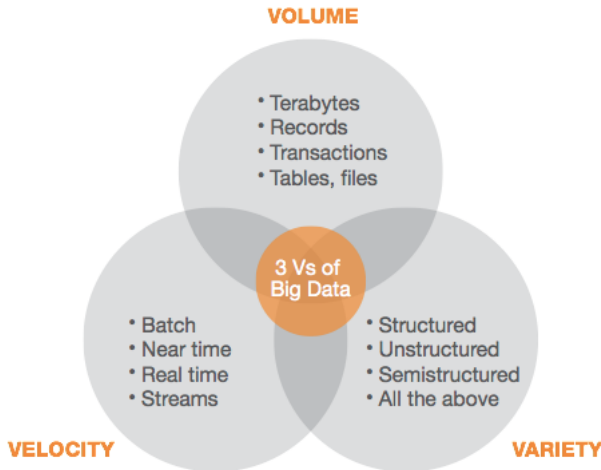
- 2.5 quintillion ( $1 \times 10^{18}$ ) bytes of data/day
- 90% generated in last 2 year. 80% of it is unstructured.
- Facebook, 4 million posts/min
- Aircraft: each engine generates 10 TB data in 30 min. UK to New York would generate 640 TB of data



<sup>3</sup><http://wersm.com/how-much-data-is-generated-every-minute-on-social-media/>

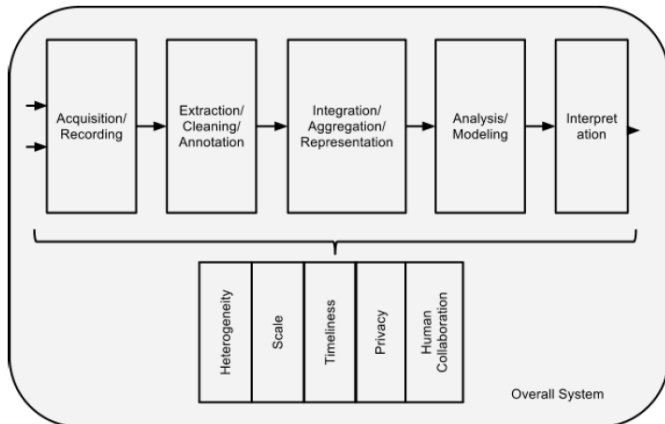
# Big Data

3V's: The “BIG” in big data is not only for volume



# Big Data

Major steps in analysis of big data: <sup>4</sup>



Solution is **Distributed computing**

<sup>4</sup>“Challenges and Opportunities with Big Data” A community white paper developed by leading researchers across the United States



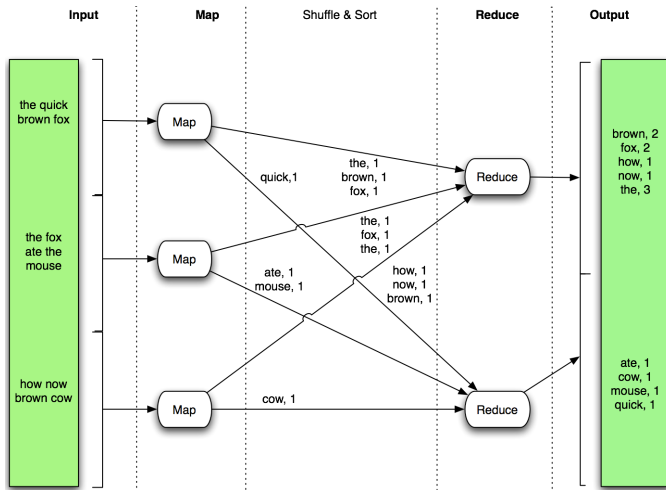
# Big Data

- **Grid** (heterogeneous) or **cluster** (homogeneous) computing
- Distributed computing **have to deal with** synchronization, deadlocks, data dependency, mutual exclusion, replication, reliability, platform scalability and provisioning
- Ready-made solution is **MapReduce/Hadoop**<sup>5</sup> or **spark** that provides a high level of abstraction for data parallel tasks
- Hadoop/PIG combo is very effective
- Need is there for a scalable distributed computing framework that provides both abstraction and performance (by exploiting all kinds of parallelisms that exist in an algorithm)

---

<sup>5</sup>MapReduce: Simplified Data Processing on Large Clusters Jeffrey Dean and Sanjay Ghemawat, OSDI, 2004

# Big Data



**map** is a function that executes on each partition.

# PK-Means

- **K-Means:** steps involved are

- 1 Select Seed
- 2 Assignment //(most compute intensive)
- 3 Compute Centroid

- **PK-Means:** <sup>6</sup>

- ▶ Input dataset is stored on GFS/HDFS. A sequence file of  $< key, value >$  pairs. (Key: offset, Value: string of whole record)
- ▶ Dataset is split and globally, and broadcast to all mappers

## Map

- ▶ Distance calculations are parallel executed
- ▶ Each mapper has array of centers
- ▶ Computes closest center for each sample
- ▶ Intermediate values/output:  $< key, value >$  (Key: index of the closed centre, value: sample)

---

<sup>6</sup>Du, Zhihua and Wang, Yiwei and Ji, Zhen, “PK-means: A new algorithm for gene clustering” in Computational Biology and Chemistry, pages=243–247, vol 32(4), Elsevier 2008

# PK-Means

## ● PK-Means: Combiner

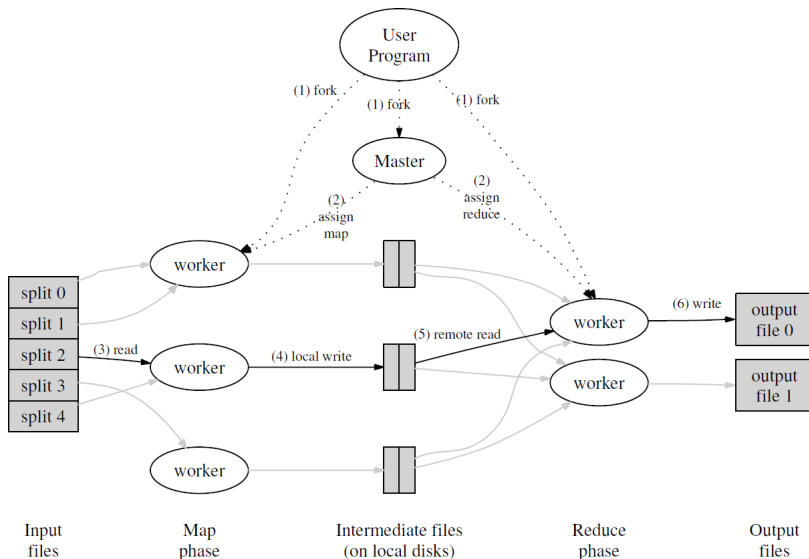
- ▶ Combines intermediate data of each map task and stores locally
- ▶ Partial sum the values assigned to the same cluster
  - ★ Record number of samples in each cluster and
  - ★ Sum of values at each dimension
- ▶ Key: key & Value: string of num and sums

## Reducer

- ▶ Input: output of combiner
- ▶ Compute all the samples assigned to a center
- ▶ Calculate new centers

**Scalability** is high

# How MapReduce Works




# How MapReduce Works

- Map function invokes data partitioning across multiple machines
  - ▶ M map tasks and R reduce tasks
  - ▶ A map worker reads the corresponding input, parses key value pair and execute user defined map function
  - ▶ Intermediate key value pair is generated
  - ▶ Periodically buffered pairs are written to local disk, partitioned into R regions and locations of the partitions are passed to the master
- Reduce invocations partition intermediate key result into R pieces
  - ▶ Reduce worker is notified by the master about partitions
  - ▶ Remote procedure calls are used to read partitioned data
  - ▶ After reading data is sorted by reduce worker and user defined reduce function is execute
  - ▶ The output of the reduce function is appended in a file
- The number of partitions (R) and the partitioning function are specified by the user

# MR-DBSCAN


- MR-DBSCAN<sup>7</sup> involves following steps
  - 1 Preprocessing
  - 2 Local DBSCAN
  - 3 Find Merging Mapping
- Uses quadtree, a spacial data structure
- Extended regions ( $\epsilon$ -extended) is taken in partition




<sup>7</sup>MR-DBSCAN: An Efficient Parallel Density-based Clustering Algorithm using MapReduce 2011 IEEE 17th International Conference on Parallel and Distributed Systems

# MR-DBSCAN

- Extended regions ( $\epsilon$ -extended) is taken in partition



- Cross connection files are processed during reduce
- This makes the algorithm data parallel





# Thank You!

**Thank you very much for your attention!**

**Queries ?**